

DYNAMO: Towards Automated Network Attack Attribution via Density-Aware Active Learning

19/01/2024

Hélène Orsini, Yufei Han

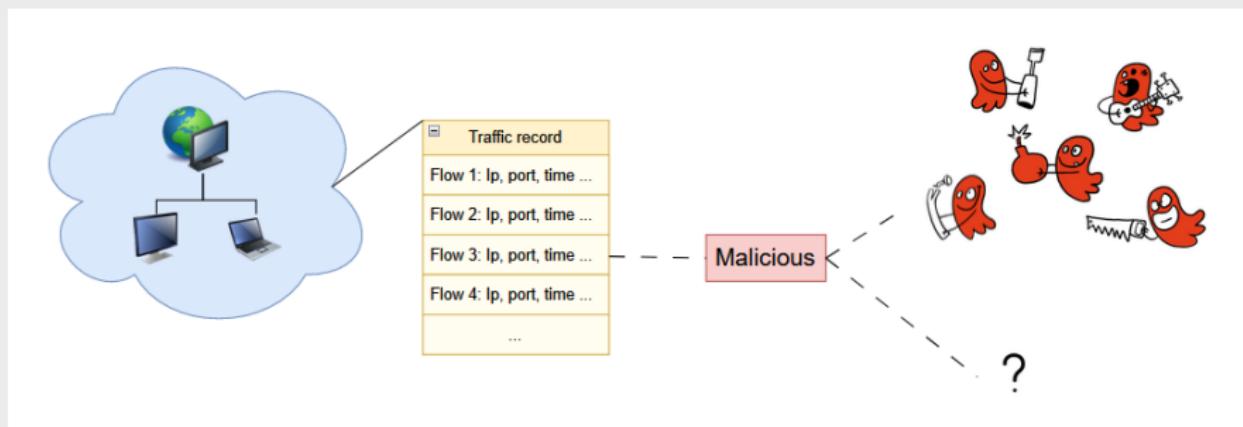
1. Introduction

2. DYNAMO Design

3. Experiment

4. Conclusion

Introduction



Introduction

Focus

- Network traffic attribution
- Machine learning techniques

Traditional challenges

- **Annotation effort**
- **Imbalance in data volumes**
- Ever evolving traffic (concept drift)

Our Proposed Framework: DYNAMO

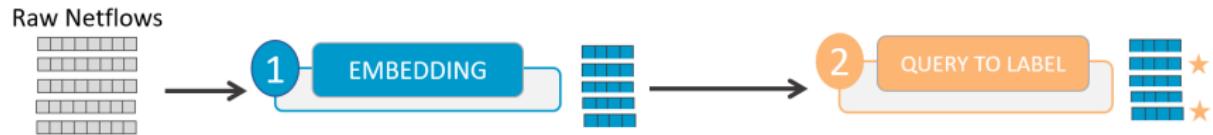
Raw Netflows



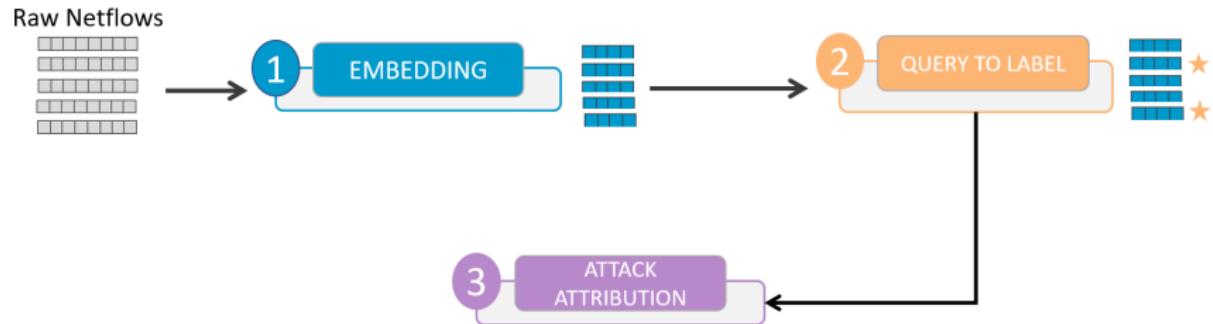
Our Proposed Framework: DYNAMO



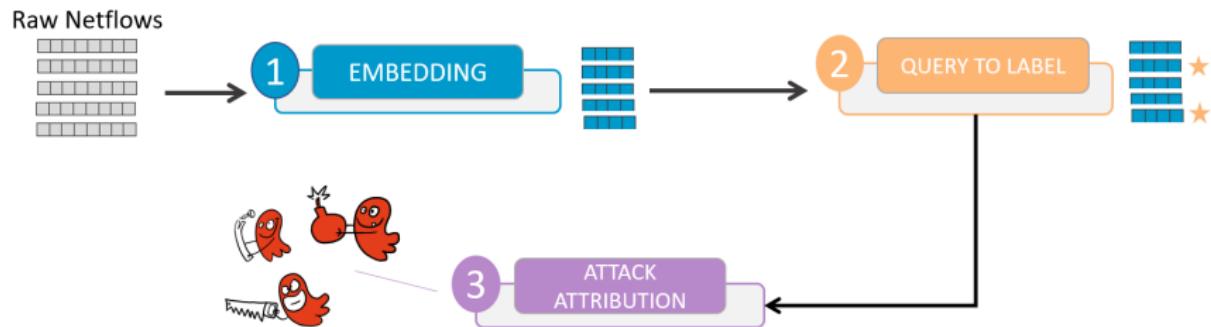
Our Proposed Framework: DYNAMO



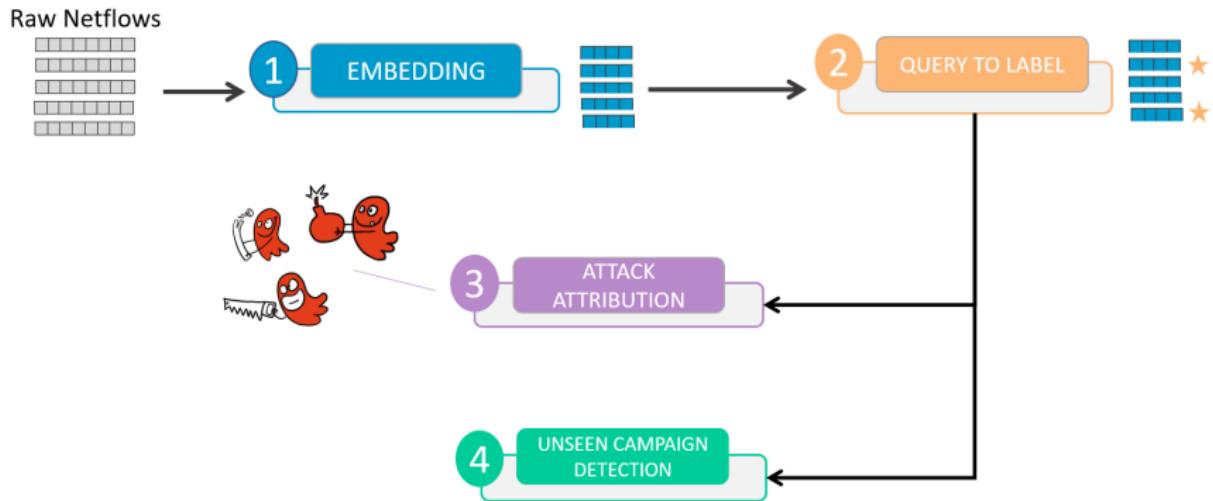
Our Proposed Framework: DYNAMO



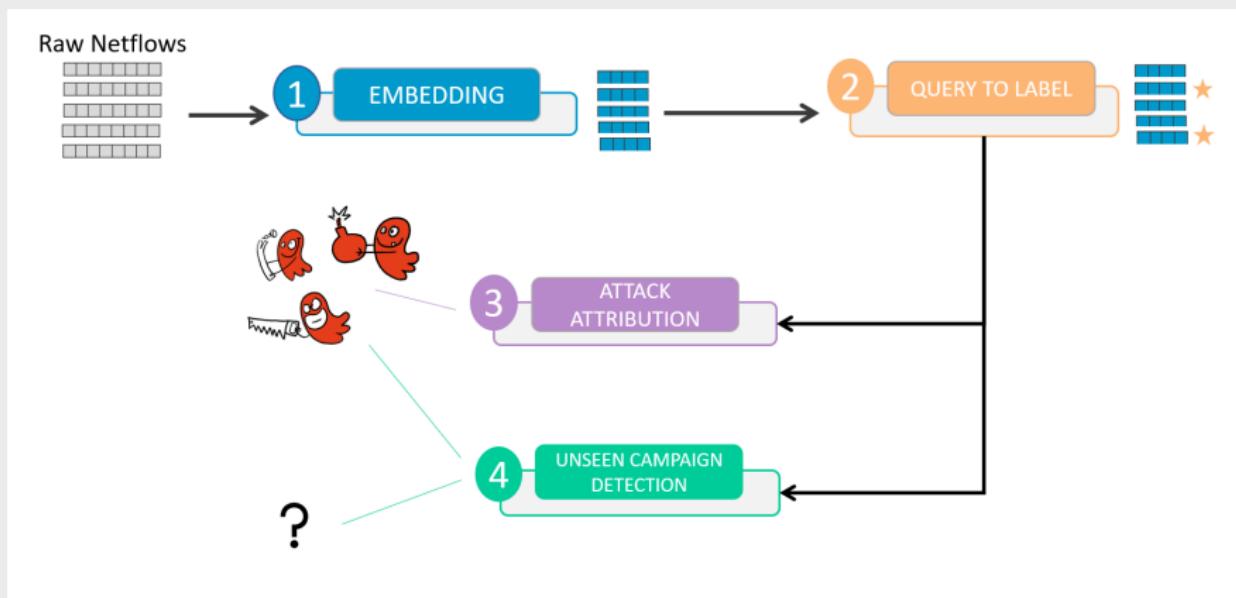
Our Proposed Framework: DYNAMO



Our Proposed Framework: DYNAMO



Our Proposed Framework: DYNAMO



Related work

Attack attribution

- Manual analysis: synthesizing and analyze report
- Machine learning-based: multi-class classification task

Query to Label: Active learning

- Uncertainty-based sampling
- Representation-based sampling

1. Introduction
2. DYNAMO Design
3. Experiment
4. Conclusion

Method

1

EMBEDDING

2

QUERY TO LABEL

3

ATTACK
ATTRIBUTION

4

UNSEEN CAMPAIGN
DETECTION

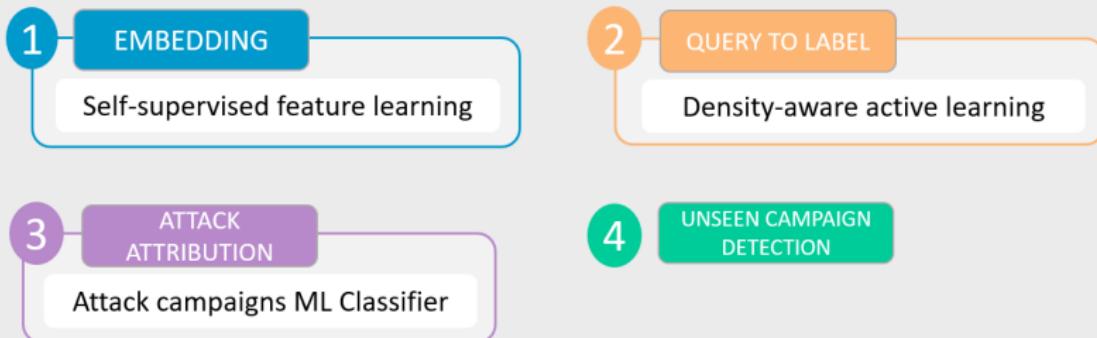
Method



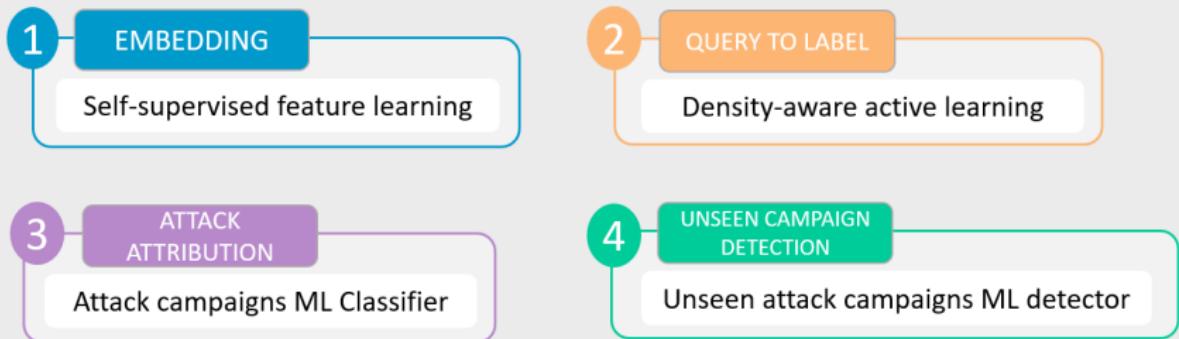
Method



Method



Method



Nearest neighbor-based self-supervised feature encoding

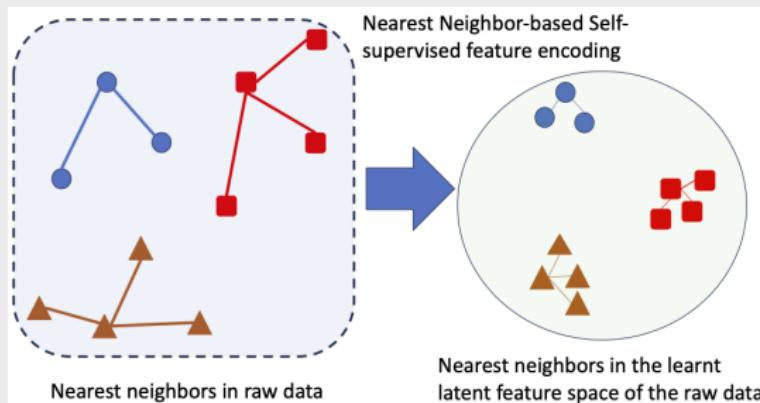
1

EMBEDDING

Raw feature vectors from Netflow : GraphSage method

$\theta^* =$

$$\arg \min_{\theta} -\frac{1}{nK} \sum_{i=1}^n [\sum_{k=1}^K \log(\sigma(h_{\theta}^T(x_i) h_{\theta}(x_i^{NN,k}))) - \lambda \sum_{j, x_v \notin KNN(x_i)} \log(\sigma(-h_{\theta}^T(x_i) h_{\theta}(x_v)))]$$

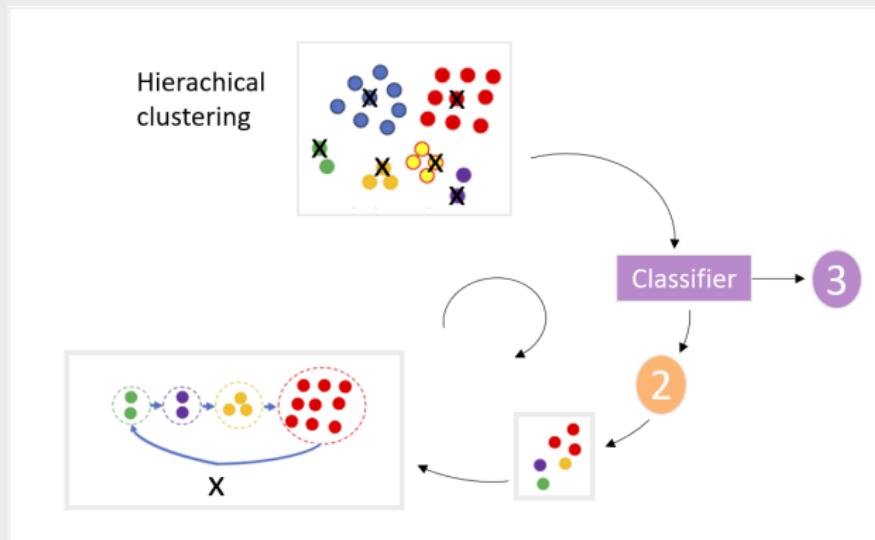


Density-aware active learning

2

QUERY TO LABEL

3

ATTACK
ATTRIBUTION

Unseen campaign strategy

4

UNSEEN CAMPAIGN
DETECTION

Pu learning

Train a classifier to distinguish between positive and negative.

Learning phase: **Positive and Unlabelled** (*only some of the positive examples in the training data are labeled and none of the negative examples are*)

$$g_{\phi}^{pu} = \arg \min_{\phi} \frac{\pi}{n_p} \sum_{x_i \in S} [\ell(g_{\phi}^{pu}(h_{\theta}(x_i)), y_i = +1) - \ell(g_{\phi}^{pu}(h_{\theta}(x_i)), y_i = -1)] + \frac{1}{n_u} \sum_{x_i \in X_{unlabeled}} \ell(g_{\phi}^{pu}(h_{\theta}(x_i)), y_i = -1)$$

1. Introduction

2. DYNAMO Design

3. Experiment

4. Conclusion

Goals

Q1 - Raw vs Embedded

Q2 - Effectiveness of DYNAMO's density-aware active learning module ?

Q3 - Effectiveness ML-based unseen campaign detection ?

Set up

Dataset: CTU13

D_{attr}^{test}	D_{ood}^{test}
D_{attr}^{train}	D_{ood}^{train}

Scenario	Flows	%
1	39933	9,23
2	18 8839	4,35
3	26 759	6,18
4	1 719	0,4
5	695	0,16
6	4 431	1,02
7	37	0,0085
8	5 052	1,17
9	179 880	41,57
10	106 315	24,57
11	8 161	1,89
12	2 143	0,50
13	38 791	8,96

Set up

Dataset: CTU13

D_{attr}^{test}	D_{ood}^{test}
D_{attr}^{train}	D_{ood}^{train}

Scenario	Flows	%
1	39933	9,23
2	18 8839	4,35
3	26 759	6,18
4	1 719	0,4
5	695	0,16
6	4 431	1,02
7	37	0,0085
8	5 052	1,17
9	179 880	41,57
10	106 315	24,57
11	8 161	1,89
12	2 143	0,50
13	38 791	8,96

Baseline 2 QUERY TO LABEL

- Select $p\%$ of D_{attr}^{train} (1k, 2k, 3k, 4k, and 5k)
- Random, UAL, and DYNAMO

Set up

Dataset: CTU13

D_{attr}^{test}	D_{ood}^{test}
D_{attr}^{train}	D_{ood}^{train}

Scenario	Flows	%
1	39933	9,23
2	18 8839	4,35
3	26 759	6,18
4	1 719	0,4
5	695	0,16
6	4 431	1,02
7	37	0,0085
8	5 052	1,17
9	179 880	41,57
10	106 315	24,57
11	8 161	1,89
12	2 143	0,50
13	38 791	8,96

Baseline 2 QUERY TO LABEL

- Select $p\%$ of D_{attr}^{train} (1k, 2k, 3k, 4k, and 5k)
- Random, UAL, and DYNAMO

Attack attribution 3 ATTACK ATTRIBUTION

- Gradient Boosting Trees, Label Spreading
- Macro F1, Balanced Accuracy

Set up

Dataset: CTU13

D_{attr}^{test}	D_{ood}^{test}
D_{attr}^{train}	D_{ood}^{train}

Scenario	Flows	%
1	39933	9,23
2	18 8839	4,35
3	26 759	6,18
4	1 719	0,4
5	695	0,16
6	4 431	1,02
7	37	0,0085
8	5 052	1,17
9	179 880	41,57
10	106 315	24,57
11	8 161	1,89
12	2 143	0,50
13	38 791	8,96

Baseline

2

QUERY TO LABEL

- Select $p\%$ of D_{attr}^{train} (1k, 2k, 3k, 4k, and 5k)
- Random, UAL, and DYNAMO

Attack attribution

3

ATTACK ATTRIBUTION

- Gradient Boosting Trees, Label Spreading
- Macro F1, Balanced Accuracy

Unseen campaign detection

4

UNSEEN CAMPAIGN DETECTION

- ISO, OCSVM, and PU
- Macro F1, AUC

Results - Attack attribution Q1

Mean \pm Standard deviation of GBT trained with the full supervision method

NB	Raw		Embedding	
	Macro F1	Balanced Acc	Macro F1	Balanced Acc
29,918 (p=20%)	0.644 \pm 0.012	0.637 \pm 0.015	0.770 \pm 0.011	0.74 \pm 0.015
59,835 (p=40%)	0.687 \pm 0.009	0.648 \pm 0.010	0.780 \pm 0.006	0.754 \pm 0.007
89,754, (p=60%)	0.675 \pm 0.006	0.637 \pm 0.006	0.792 \pm 0.004	0.765 \pm 0.005
119,671 (p=80%)	0.685 \pm 0.006	0.665 \pm 0.006	0.801 \pm 0.002	0.772 \pm 0.003
149,589 (p=100%)	0.674 \pm 0.006	0.670 \pm 0.002	0.805 \pm 0.002	0.786 \pm 0.003



Embedded data increases both macro-F1 score and balanced accuracy.

Results - Attack attribution Q2

Mean \pm Standard deviation of Macro F1-score

		Attack attribution with the latent feature learned by the self-supervised learning module					
		Random Selection		DYNAMO		UAL	
NB		GB	LS	GB	LS	GB	LS
1000 ($p=0.7\%$)		0.611 \pm 0.024	0.637 \pm 0.036	0.695 \pm 0.024	0.631 \pm 0.000	0.607 \pm 0.016	0.574 \pm 0.067
2000 ($p=1.3\%$)		0.653 \pm 0.018	0.694 \pm 0.022	0.745 \pm 0.021	0.677 \pm 0.000	0.613 \pm 0.016	0.608 \pm 0.017
3000 ($p=2.0\%$)		0.673 \pm 0.017	0.712 \pm 0.016	0.764 \pm 0.016	0.688 \pm 0.000	0.723 \pm 0.013	0.654 \pm 0.027
4000 ($p=2.6\%$)		0.686 \pm 0.013	0.723 \pm 0.049	0.781 \pm 0.015	0.707 \pm 0.000	0.773 \pm 0.002	0.689 \pm 0.019
5000 ($p=3.3\%$)		0.697 \pm 0.013	0.732 \pm 0.012	0.791 \pm 0.011	0.708 \pm 0.000	0.785 \pm 0.009	0.702 \pm 0.020



Dynamo outperforms UAL (GB, LS) and Random (GB)

Results - Unseen campaign Q3

Mean \pm Standard deviation of Macro F1-score

Unseen campaign detection with the latent feature learned by the self-supervised learning module									
	Random Selection			DYNAMO			UAL		
NB	ISO	OCSVM	PU	ISO	OCSVM	PU	ISO	OCSVM	PU
1000 ($p=0.7\%$)	0.748 \pm 0.005	0.853 \pm 0.000	1.000 \pm 0.000	0.913 \pm 0.007	0.921 \pm 0.005	1.000 \pm 0.000	0.832 \pm 0.043	0.898 \pm 0.013	1.000 \pm 0.000
2000 ($p=1.3\%$)	0.754 \pm 0.005	0.762 \pm 0.000	1.000 \pm 0.000	0.905 \pm 0.010	0.913 \pm 0.006	1.000 \pm 0.000	0.817 \pm 0.044	0.880 \pm 0.016	1.000 \pm 0.000
3000 ($p=2.0\%$)	0.672 \pm 0.005	0.696 \pm 0.000	1.000 \pm 0.000	0.765 \pm 0.009	0.909 \pm 0.008	1.000 \pm 0.000	0.789 \pm 0.018	0.860 \pm 0.021	1.000 \pm 0.000
4000 ($p=2.6\%$)	0.758 \pm 0.007	0.687 \pm 0.000	1.000 \pm 0.000	0.897 \pm 0.010	0.904 \pm 0.010	1.000 \pm 0.000	0.789 \pm 0.046	0.848 \pm 0.048	1.000 \pm 0.000
5000 ($p=3.3\%$)	0.754 \pm 0.007	0.689 \pm 0.026	1.000 \pm 0.000	0.891 \pm 0.009	0.898 \pm 0.008	1.000 \pm 0.000	0.794 \pm 0.059	0.842 \pm 0.068	1.000 \pm 0.000



Pu performs best

Dynamo provide better result for ISO and OCSVM

1. Introduction

2. DYNAMO Design

3. Experiment

4. Conclusion

Key Takeaway

- ML-Based attack attribution challenges: scarce label data, imbalanced campaign distribution
- Self-Supervised feature encoding boosts attack attribution
- Density-aware active learning helps overcome imbalanced data issue
- Positive-Unlabeled learning outperforms for unseen campaign detection

Summary and perspectives

- DYNAMO, a weakly supervised ML-based pipeline designed for automated network attack attribution without requiring exhaustive labeling of attack campaigns
- Density-aware Machine Learning
- Positive-unlabeled weakly supervised learning

Next steps

- Adapt to new dataset (Poneypot)
- Add transfer learning environment

Thanks, Questions ?

